

simclr_mnist

December 13, 2023

1 Contrastive Learning on non-spurious MNIST dataset

```
[1]: import torch
import torch.optim as optim
import torch.nn as nn
import torch.nn.functional as F

from torchvision import datasets, transforms
import matplotlib.pyplot as plt
import numpy as np
```

```
[35]: torch.cuda.is_available()
```

```
[35]: True
```

2 Utils

```
[2]: from collections import defaultdict

class MetricMonitor:
    def __init__(self, float_precision=4):
        self.float_precision = float_precision
        self.reset()

    def reset(self):
        self.metrics = defaultdict(lambda: {"val": 0, "count": 0, "avg": 0})

    def update(self, metric_name, val):
        metric = self.metrics[metric_name]

        metric["val"] += val
        metric["count"] += 1
        metric["avg"] = metric["val"] / metric["count"]

    def __str__(self):
        return " | ".join(
```

```

        [
            "{metric_name}: {avg:.{float_precision}f}".format(
                metric_name=metric_name, avg=metric["avg"],
                float_precision=self.float_precision
            )
            for (metric_name, metric) in self.metrics.items()
        ]
    )

```

```

[3]: class EarlyStopping:
    def __init__(self, patience=7, verbose=False, delta=0, path='checkpoint.
    pt', trace_func=print):
        """
        Source:
            https://github.com/Bjarten/early-stopping-pytorch/blob/master/
            pytorchtools.py
        Args:
            patience (int): How long to wait after last time validation loss
            improved.
                                Default: 7
            verbose (bool): If True, prints a message for each validation loss
            improvement.
                                Default: False
            delta (float): Minimum change in the monitored quantity to qualify
            as an improvement.
                                Default: 0
            path (str): Path for the checkpoint to be saved to.
                                Default: 'checkpoint.pt'
            trace_func (function): trace print function.
                                Default: print
        """
        self.patience = patience
        self.verbose = verbose
        self.counter = 0
        self.best_score = None
        self.early_stop = False
        self.val_loss_min = np.Inf
        self.delta = delta
        self.path = path
        self.trace_func = trace_func
    def __call__(self, val_loss, model):

        score = -val_loss

        if self.best_score is None:
            self.best_score = score

```

```

        self.save_checkpoint(val_loss, model)
    elif score < self.best_score + self.delta:
        self.counter += 1
        self.trace_func(f'> early stopping counter: {self.counter} out of_{self.patience}')
        if self.counter >= self.patience:
            self.early_stop = True
        else:
            self.best_score = score
            self.save_checkpoint(val_loss, model)
            self.counter = 0

    def save_checkpoint(self, val_loss, model):
        '''Saves model when validation loss decrease.'''
        if self.verbose:
            self.trace_func(f'validation loss decreased ({self.val_loss_min:.6f} --> {val_loss:.6f}). Saving model ...')
            torch.save(model.state_dict(), self.path)
            self.val_loss_min = val_loss

```

```

[4]: def calculate_accuracy(output, target):
    "Calculates accuracy"
    output = output.data.max(dim=1, keepdim=True)[1]
    output = output == 1.0
    output = torch.flatten(output)
    target = target == 1.0
    target = torch.flatten(target)
    return torch.true_divide((target == output).sum(dim=0), output.size(0)).item()

```

```

[5]: def save_model(model, optimizer, epoch, save_file):
    print('\n==> Saving...')
    state = {
        'model': model.state_dict(),
        'optimizer': optimizer.state_dict(),
        'epoch': epoch,
    }
    torch.save(state, save_file)
    del state

```

```

[6]: class TwoCropTransform:
    """Create two crops of the same image"""
    def __init__(self, transform):
        self.transform = transform

    def __call__(self, x):
        return [self.transform(x), self.transform(x)]

```

3 Method

```
[7]: class Encoder(torch.nn.Module):
    """Encoder network"""
    def __init__(self):
        super(Encoder, self).__init__()
        # L1 (?, 28, 28, 3) -> (?, 28, 28, 32) -> (?, 14, 14, 32)
        self.layer1 = torch.nn.Sequential(
            torch.nn.Conv2d(3, 32, kernel_size=3, stride=1, padding=1),
            torch.nn.BatchNorm2d(32),
            torch.nn.ReLU(),
            torch.nn.MaxPool2d(kernel_size=2, stride=2),
            torch.nn.Dropout(p=0.2)
        )
        # L2 (?, 14, 14, 32) -> (?, 14, 14, 64) -> (?, 7, 7, 64)
        self.layer2 = torch.nn.Sequential(
            torch.nn.Conv2d(32, 64, kernel_size=3, stride=1, padding=1),
            torch.nn.BatchNorm2d(64),
            torch.nn.ReLU(),
            torch.nn.MaxPool2d(kernel_size=2, stride=2),
            torch.nn.Dropout(p=0.2)
        )
        # L3 (?, 7, 7, 64) -> (?, 7, 7, 128) -> (?, 4, 4, 128)
        self.layer3 = torch.nn.Sequential(
            torch.nn.Conv2d(64, 128, kernel_size=3, stride=1, padding=1),
            torch.nn.BatchNorm2d(128),
            torch.nn.ReLU(),
            torch.nn.MaxPool2d(kernel_size=2, stride=2, padding=1),
            torch.nn.Dropout(p=0.2)
        )
        self._to_linear = 4 * 4 * 128

    def forward(self, x):
        x = self.layer1(x)
        x = self.layer2(x)
        x = self.layer3(x)
        x = x.view(x.size(0), -1) # Flatten them for FC
        return x

class LinearClassifier(torch.nn.Module):
    """Linear classifier"""
    def __init__(self):
        super(LinearClassifier, self).__init__()
        self.fc = torch.nn.Sequential(
            torch.nn.Linear(4 * 4 * 128, 10),
        )
```

```

def forward(self, x):
    x = self.fc(x)
    probs = torch.nn.functional.softmax(x, dim=0)
    return probs

```

```

[8]: class SupCon(nn.Module):
    """encoder + projection head"""
    def __init__(self, model, head='mlp', feat_dim=128):
        super(SupCon, self).__init__()

        self.dim_in = model._to_linear
        self.encoder = model

        if head == 'linear':
            self.head = nn.Linear(self.dim_in, feat_dim)
        elif head == 'mlp':
            self.head = nn.Sequential(
                nn.Linear(self.dim_in, self.dim_in),
                nn.ReLU(inplace=True),
                nn.Linear(self.dim_in, feat_dim)
            )
        else:
            raise NotImplementedError('Head not supported: {}'.format(head))

    def forward(self, x):
        feat = self.encoder(x)
        feat = F.normalize(self.head(feat), dim=1)
        return feat

```

```

[9]: """
    Author: Yonglong Tian (yonglong@mit.edu)
    Date: May 07, 2020
    """
    class SupConLoss(nn.Module):
        def __init__(self, temperature=0.07, contrast_mode='all',
            base_temperature=0.07):
            super(SupConLoss, self).__init__()
            self.temperature = temperature
            self.contrast_mode = contrast_mode
            self.base_temperature = base_temperature

        def forward(self, features, labels=None, mask=None):
            device = (torch.device('cuda')
                if features.is_cuda
                else torch.device('cpu'))

```

```

if len(features.shape) < 3:
    raise ValueError('`features` needs to be [bsz, n_views, ...],\n'
                     'at least 3 dimensions are required')
if len(features.shape) > 3:
    features = features.view(features.shape[0], features.shape[1], -1)

batch_size = features.shape[0]
if labels is not None and mask is not None:
    raise ValueError('Cannot define both `labels` and `mask`')
elif labels is None and mask is None:
    mask = torch.eye(batch_size, dtype=torch.float32).to(device)
elif labels is not None:
    labels = labels.contiguous().view(-1, 1)
    if labels.shape[0] != batch_size:
        raise ValueError('Num of labels does not match num of features')
    mask = torch.eq(labels, labels.T).float().to(device)
else:
    mask = mask.float().to(device)

contrast_count = features.shape[1]
contrast_feature = torch.cat(torch.unbind(features, dim=1), dim=0)
if self.contrast_mode == 'one':
    anchor_feature = features[:, 0]
    anchor_count = 1
elif self.contrast_mode == 'all':
    anchor_feature = contrast_feature
    anchor_count = contrast_count
else:
    raise ValueError('Unknown mode: {}'.format(self.contrast_mode))

# compute logits
anchor_dot_contrast = torch.div(
    torch.matmul(anchor_feature, contrast_feature.T),
    self.temperature)
# for numerical stability
logits_max, _ = torch.max(anchor_dot_contrast, dim=1, keepdim=True)
logits = anchor_dot_contrast - logits_max.detach()

# tile mask
mask = mask.repeat(anchor_count, contrast_count)
# mask-out self-contrast cases
logits_mask = torch.scatter(
    torch.ones_like(mask),
    1,
    torch.arange(batch_size * anchor_count).view(-1, 1).to(device),
    0
)

```

```

mask = mask * logits_mask

# compute log_prob
exp_logits = torch.exp(logits) * logits_mask
log_prob = logits - torch.log(exp_logits.sum(1, keepdim=True))

# compute mean of log-likelihood over positive
mean_log_prob_pos = (mask * log_prob).sum(1) / mask.sum(1)

# loss
loss = - (self.temperature / self.base_temperature) * mean_log_prob_pos
loss = loss.view(anchor_count, batch_size).mean()

return loss

```

4 Training

```

[10]: def pretraining(epoch, model, contrastive_loader, optimizer, criterion,
    ↪method='SimCLR'):
    "Contrastive pre-training over an epoch"
    metric_monitor = MetricMonitor()
    model.train()
    for batch_idx, (data, labels) in enumerate(contrastive_loader):
        # print(batch_idx)
        # print(data[0].shape)
        # print(data[1].shape)
        # print(labels.shape)
        data = torch.cat([data[0], data[1]], dim=0)
        # print(data.shape)
        if torch.cuda.is_available():
            data, labels = data.cuda(), labels.cuda()
        data, labels = torch.autograd.Variable(data, False), torch.autograd.
    ↪Variable(labels)
        bsz = labels.shape[0]
        features = model(data)
        f1, f2 = torch.split(features, [bsz, bsz], dim=0)
        features = torch.cat([f1.unsqueeze(1), f2.unsqueeze(1)], dim=1)
        if method == 'SupCon':
            loss = criterion(features, labels)
        elif method == 'SimCLR':
            loss = criterion(features)
        else:
            raise ValueError('contrastive method not supported: {}'.
    ↪format(method))
        metric_monitor.update("Loss", loss.item())
        metric_monitor.update("Learning Rate", optimizer.param_groups[0]['lr'])

```

```

        optimizer.zero_grad()
        loss.backward()
        optimizer.step()
        print("[Epoch: {epoch:03d}] Contrastive Pre-train | {metric_monitor}").
        ↪format(epoch=epoch, metric_monitor=metric_monitor))
        return metric_monitor.metrics['Loss']['avg'], metric_monitor.
        ↪metrics['Learning Rate']['avg']

```

```

[11]: def training(epoch, model, classifier, train_loader, optimizer, criterion):
        "Training over an epoch"
        metric_monitor = MetricMonitor()
        model.eval()
        classifier.train()
        for batch_idx, (data, labels) in enumerate(train_loader):
            # print(batch_idx)
            # print(data.shape)
            # print(labels.shape)
            if torch.cuda.is_available():
                data, labels = data.cuda(), labels.cuda()
            data, labels = torch.autograd.Variable(data, False), torch.autograd.
            ↪Variable(labels)
            with torch.no_grad():
                features = model.encoder(data)
                output = classifier(features.float())
                loss = criterion(output, labels)
                accuracy = calculate_accuracy(output, labels)
                metric_monitor.update("Loss", loss.item())
                metric_monitor.update("Accuracy", accuracy)
                data.detach()
                labels.detach()
                optimizer.zero_grad()
                loss.backward()
                optimizer.step()
            print("[Epoch: {epoch:03d}] Train      | {metric_monitor}").
            ↪format(epoch=epoch, metric_monitor=metric_monitor))
            return metric_monitor.metrics['Loss']['avg'], metric_monitor.
            ↪metrics['Accuracy']['avg']

```

```

[12]: def validation(epoch, model, classifier, valid_loader, criterion):
        "Validation over an epoch"
        metric_monitor = MetricMonitor()
        model.eval()
        classifier.eval()
        with torch.no_grad():
            for batch_idx, (data, labels) in enumerate(valid_loader):
                if torch.cuda.is_available():
                    data, labels = data.cuda(), labels.cuda()

```



```

        data, labels = torch.autograd.Variable(data, False), torch.autograd.
↪Variable(labels)
        features = model.encoder(data)
        output = classifier(features.float())
        loss = criterion(output, labels)
        accuracy = calculate_accuracy(output, labels)
        metric_monitor.update("Loss", loss.item())
        metric_monitor.update("Accuracy", accuracy)
        data.detach()
        labels.detach()

        print("[Epoch: {epoch:03d}] Validation | {metric_monitor}".
↪format(epoch=epoch, metric_monitor=metric_monitor))
        return metric_monitor.metrics['Loss']['avg'], metric_monitor.
↪metrics['Accuracy']['avg']

```

5 Data

```

[13]: contrastive_transform = transforms.Compose([
    transforms.Grayscale(num_output_channels = 3), # Convert to 3 color channel
    transforms.RandomHorizontalFlip(),
    transforms.RandomResizedCrop(size=28, scale=(0.2, 1.)),
    transforms.ToTensor(),
    transforms.Normalize((0.5, ), (0.5, )),
])
train_transform = transforms.Compose([
    transforms.Grayscale(num_output_channels = 3), # Convert to 3 color channel
    transforms.RandomHorizontalFlip(),
    transforms.ToTensor(),
    transforms.Normalize((0.5, ), (0.5, )),
])
valid_transform = transforms.Compose([
    transforms.Grayscale(num_output_channels = 3), # Convert to 3 color channel
    transforms.ToTensor(),
    transforms.Normalize((0.5, ), (0.5, )),
])

contrastive_set = datasets.MNIST('./data', download=True, train=True,
↪transform=TwoCropTransform(contrastive_transform))
train_set = datasets.MNIST('./data', download=True, train=True,
↪transform=train_transform)
valid_set = datasets.MNIST('./data', download=True, train=False,
↪transform=valid_transform)

```

Downloading <http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz>
 Downloading <http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz> to
 ./data/MNIST/raw/train-images-idx3-ubyte.gz

```

100%|      | 9912422/9912422 [00:00<00:00, 176856960.98it/s]
Extracting ./data/MNIST/raw/train-images-idx3-ubyte.gz to ./data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz to
./data/MNIST/raw/train-labels-idx1-ubyte.gz
100%|      | 28881/28881 [00:00<00:00, 33124335.20it/s]
Extracting ./data/MNIST/raw/train-labels-idx1-ubyte.gz to ./data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz to
./data/MNIST/raw/t10k-images-idx3-ubyte.gz
100%|      | 1648877/1648877 [00:00<00:00, 39444547.47it/s]
Extracting ./data/MNIST/raw/t10k-images-idx3-ubyte.gz to ./data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz to
./data/MNIST/raw/t10k-labels-idx1-ubyte.gz
100%|      | 4542/4542 [00:00<00:00, 5719162.04it/s]
Extracting ./data/MNIST/raw/t10k-labels-idx1-ubyte.gz to ./data/MNIST/raw

```

```

[14]: print(contrastive_set)
      print(contrastive_set.classes)
      print(len(contrastive_set.classes))
      print(len(contrastive_set))
      print(len(contrastive_set[0])) # 2 entries (2 crops of same img for a given
      ↪ label)
      print(len(contrastive_set[0][0])) # img, label
      print(len(contrastive_set[0][0][0])) # color channel
      print(len(contrastive_set[0][0][0][0])) # dim
      print(len(contrastive_set[0][0][0][0][0])) # dim

```

Dataset MNIST

```

Number of datapoints: 60000
Root location: ./data
Split: Train
StandardTransform

```

Transform: <__main__.TwoCropTransform object at 0x7913f5092a70>

```

['0 - zero', '1 - one', '2 - two', '3 - three', '4 - four', '5 - five', '6 -
six', '7 - seven', '8 - eight', '9 - nine']

```

```
10
60000
2
2
3
28
28
```

```
[15]: print(train_set)
      print(train_set.classes)
      print(len(train_set.classes))
      print(len(train_set))
      print(len(train_set[0])) # img, label
      print(len(train_set[0][0])) # color channel
      print(len(train_set[0][0][0])) # dim
      print(len(train_set[0][0][0][0])) # dim
```

Dataset MNIST

Number of datapoints: 60000

Root location: ./data

Split: Train

StandardTransform

Transform: Compose(

Grayscale(num_output_channels=3)

RandomHorizontalFlip(p=0.5)

ToTensor()

Normalize(mean=(0.5,), std=(0.5,))

)

['0 - zero', '1 - one', '2 - two', '3 - three', '4 - four', '5 - five', '6 - six', '7 - seven', '8 - eight', '9 - nine']

```
10
60000
2
3
28
28
```

```
[16]: print(valid_set)
      print(valid_set.classes)
      print(len(valid_set.classes))
      print(len(valid_set))
      print(len(valid_set[0])) # img, label
      print(len(valid_set[0][0])) # color channel
      print(len(valid_set[0][0][0])) # dim
      print(len(valid_set[0][0][0][0])) # dim
```

Dataset MNIST

Number of datapoints: 10000

```

Root location: ./data
Split: Test
StandardTransform
Transform: Compose(
    Grayscale(num_output_channels=3)
    ToTensor()
    Normalize(mean=(0.5,), std=(0.5,))
)
['0 - zero', '1 - one', '2 - two', '3 - three', '4 - four', '5 - five', '6 -
six', '7 - seven', '8 - eight', '9 - nine']
10
10000
2
3
28
28

```

```

[33]: import torchvision.transforms as T
      T.ToPILImage()(contrastive_set[0][0][0]).resize((50,50))

```

[33]:



6 4. Main

```

[36]: import os

def main():

    num_epochs = 50
    use_early_stopping = True
    use_scheduler = True
    head_type = 'mlp' # choose among 'mlp' and 'linear'
    save_file = os.path.join('./results/', 'model.pth')
    if not os.path.isdir('./results/'):
        os.makedirs('./results/')

    contrastive_loader = torch.utils.data.DataLoader(contrastive_set,
↳ batch_size=64, shuffle=True)
    train_loader = torch.utils.data.DataLoader(train_set, batch_size=64,
↳ shuffle=True)
    valid_loader = torch.utils.data.DataLoader(valid_set, batch_size=64,
↳ shuffle=True)

```

```

# Part 1
encoder = Encoder()
model = SupCon(encoder, head=head_type, feat_dim=128)
criterion = SupConLoss(temperature=0.07)
if torch.cuda.is_available():
    model = model.cuda()
    criterion = criterion.cuda()
optimizer = torch.optim.Adam(model.parameters(), lr=1e-3, betas=(0.9, 0.
↪999), eps=1e-08, weight_decay=1e-3)
scheduler = torch.optim.lr_scheduler.StepLR(optimizer, step_size=20,
↪gamma=0.9)

contrastive_loss, contrastive_lr = [], []

for epoch in range(1, num_epochs+1):
    print('Epoch {} start'.format(epoch))
    loss, lr = pretraining(epoch, model, contrastive_loader, optimizer,
↪criterion, method='SimCLR')
    if use_scheduler:
        scheduler.step()
    contrastive_loss.append(loss)
    contrastive_lr.append(lr)
    print('Epoch {} end'.format(epoch))

save_model(model, optimizer, num_epochs, save_file)

plt.plot(range(1, len(contrastive_lr)+1), contrastive_lr, color='b', label =
↪'learning rate')
plt.legend(), plt.ylabel('loss'), plt.xlabel('epochs'), plt.title('Learning
↪Rate'), plt.show()

plt.plot(range(1, len(contrastive_loss)+1), contrastive_loss, color='b',
↪label = 'loss')
plt.legend(), plt.ylabel('loss'), plt.xlabel('epochs'), plt.title('Loss'),
↪plt.show()

# Part 2
model = SupCon(encoder, head=head_type, feat_dim=128)
classifier = LinearClassifier()
criterion = torch.nn.CrossEntropyLoss()

ckpt = torch.load(save_file, map_location='cpu')
state_dict = ckpt['model']
new_state_dict = {}
for k, v in state_dict.items():

```

```

        k = k.replace("module.", "")
        new_state_dict[k] = v
    state_dict = new_state_dict
    model.load_state_dict(state_dict)

    if torch.cuda.is_available():
        model = model.cuda()
        classifier = classifier.cuda()
        criterion = criterion.cuda()

    train_losses , train_accuracies = [], []
    valid_losses , valid_accuracies = [], []

    if use_early_stopping:
        early_stopping = EarlyStopping(patience=30, verbose=False, delta=1e-4)

    for epoch in range(1, num_epochs+1):
        print('Epoch {} start'.format(epoch))

        train_loss, train_accuracy = training(epoch, model, classifier,
        ↪train_loader, optimizer, criterion)
        valid_loss, valid_accuracy = validation(epoch, model, classifier,
        ↪valid_loader, criterion)

        if use_scheduler:
            scheduler.step()

        train_losses.append(train_loss)
        train_accuracies.append(train_accuracy)
        valid_losses.append(valid_loss)
        valid_accuracies.append(valid_accuracy)

        print('Epoch {} end'.format(epoch))

        if use_early_stopping:
            early_stopping(valid_loss, model)

            if early_stopping.early_stop:
                print('Early stopping at epoch', epoch)
                break

    plt.plot(range(1, len(train_losses)+1), train_losses, color='b', label =
    ↪'training loss')
    plt.plot(range(1, len(valid_losses)+1), valid_losses, color='r',
    ↪linestyle='dashed', label = 'validation loss')
    plt.legend(), plt.ylabel('loss'), plt.xlabel('epochs'), plt.title('Loss'),
    ↪plt.show()

```

```

plt.plot(range(1,len(train_accuracies)+1),train_accuracies, color='b',
↪label = 'training accuracy')
plt.plot(range(1,len(valid_accuracies)+1),valid_accuracies, color='r',
↪linestyle='dashed', label = 'validation accuracy')
plt.legend(), plt.ylabel('loss'), plt.xlabel('epochs'), plt.
↪title('Accuracy'), plt.show()

```

[37]: main()

```

Epoch 1 start
[Epoch: 001] Contrastive Pre-train | Loss: 2.6896 | Learning Rate: 0.0010
Epoch 1 end
Epoch 2 start
[Epoch: 002] Contrastive Pre-train | Loss: 1.7506 | Learning Rate: 0.0010
Epoch 2 end
Epoch 3 start
[Epoch: 003] Contrastive Pre-train | Loss: 1.4733 | Learning Rate: 0.0010
Epoch 3 end
Epoch 4 start
[Epoch: 004] Contrastive Pre-train | Loss: 1.3309 | Learning Rate: 0.0010
Epoch 4 end
Epoch 5 start
[Epoch: 005] Contrastive Pre-train | Loss: 1.2423 | Learning Rate: 0.0010
Epoch 5 end
Epoch 6 start
[Epoch: 006] Contrastive Pre-train | Loss: 1.1636 | Learning Rate: 0.0010
Epoch 6 end
Epoch 7 start
[Epoch: 007] Contrastive Pre-train | Loss: 1.1083 | Learning Rate: 0.0010
Epoch 7 end
Epoch 8 start
[Epoch: 008] Contrastive Pre-train | Loss: 1.0682 | Learning Rate: 0.0010
Epoch 8 end
Epoch 9 start
[Epoch: 009] Contrastive Pre-train | Loss: 1.0319 | Learning Rate: 0.0010
Epoch 9 end
Epoch 10 start
[Epoch: 010] Contrastive Pre-train | Loss: 1.0043 | Learning Rate: 0.0010
Epoch 10 end
Epoch 11 start
[Epoch: 011] Contrastive Pre-train | Loss: 0.9945 | Learning Rate: 0.0010
Epoch 11 end
Epoch 12 start
[Epoch: 012] Contrastive Pre-train | Loss: 0.9756 | Learning Rate: 0.0010
Epoch 12 end
Epoch 13 start


```


[Epoch: 013] Contrastive Pre-train | Loss: 0.9712 | Learning Rate: 0.0010
Epoch 13 end
Epoch 14 start
[Epoch: 014] Contrastive Pre-train | Loss: 0.9434 | Learning Rate: 0.0010
Epoch 14 end
Epoch 15 start
[Epoch: 015] Contrastive Pre-train | Loss: 0.9217 | Learning Rate: 0.0010
Epoch 15 end
Epoch 16 start
[Epoch: 016] Contrastive Pre-train | Loss: 0.9292 | Learning Rate: 0.0010
Epoch 16 end
Epoch 17 start
[Epoch: 017] Contrastive Pre-train | Loss: 0.9125 | Learning Rate: 0.0010
Epoch 17 end
Epoch 18 start
[Epoch: 018] Contrastive Pre-train | Loss: 0.9121 | Learning Rate: 0.0010
Epoch 18 end
Epoch 19 start
[Epoch: 019] Contrastive Pre-train | Loss: 0.8963 | Learning Rate: 0.0010
Epoch 19 end
Epoch 20 start
[Epoch: 020] Contrastive Pre-train | Loss: 0.8896 | Learning Rate: 0.0010
Epoch 20 end
Epoch 21 start
[Epoch: 021] Contrastive Pre-train | Loss: 0.8702 | Learning Rate: 0.0009
Epoch 21 end
Epoch 22 start
[Epoch: 022] Contrastive Pre-train | Loss: 0.8589 | Learning Rate: 0.0009
Epoch 22 end
Epoch 23 start
[Epoch: 023] Contrastive Pre-train | Loss: 0.8467 | Learning Rate: 0.0009
Epoch 23 end
Epoch 24 start
[Epoch: 024] Contrastive Pre-train | Loss: 0.8478 | Learning Rate: 0.0009
Epoch 24 end
Epoch 25 start
[Epoch: 025] Contrastive Pre-train | Loss: 0.8534 | Learning Rate: 0.0009
Epoch 25 end
Epoch 26 start
[Epoch: 026] Contrastive Pre-train | Loss: 0.8387 | Learning Rate: 0.0009
Epoch 26 end
Epoch 27 start
[Epoch: 027] Contrastive Pre-train | Loss: 0.8208 | Learning Rate: 0.0009
Epoch 27 end
Epoch 28 start
[Epoch: 028] Contrastive Pre-train | Loss: 0.8265 | Learning Rate: 0.0009
Epoch 28 end
Epoch 29 start

[Epoch: 029] Contrastive Pre-train | Loss: 0.8152 | Learning Rate: 0.0009
Epoch 29 end
Epoch 30 start
[Epoch: 030] Contrastive Pre-train | Loss: 0.8211 | Learning Rate: 0.0009
Epoch 30 end
Epoch 31 start
[Epoch: 031] Contrastive Pre-train | Loss: 0.8127 | Learning Rate: 0.0009
Epoch 31 end
Epoch 32 start
[Epoch: 032] Contrastive Pre-train | Loss: 0.8060 | Learning Rate: 0.0009
Epoch 32 end
Epoch 33 start
[Epoch: 033] Contrastive Pre-train | Loss: 0.8032 | Learning Rate: 0.0009
Epoch 33 end
Epoch 34 start
[Epoch: 034] Contrastive Pre-train | Loss: 0.8028 | Learning Rate: 0.0009
Epoch 34 end
Epoch 35 start
[Epoch: 035] Contrastive Pre-train | Loss: 0.7995 | Learning Rate: 0.0009
Epoch 35 end
Epoch 36 start
[Epoch: 036] Contrastive Pre-train | Loss: 0.7959 | Learning Rate: 0.0009
Epoch 36 end
Epoch 37 start
[Epoch: 037] Contrastive Pre-train | Loss: 0.7900 | Learning Rate: 0.0009
Epoch 37 end
Epoch 38 start
[Epoch: 038] Contrastive Pre-train | Loss: 0.8028 | Learning Rate: 0.0009
Epoch 38 end
Epoch 39 start
[Epoch: 039] Contrastive Pre-train | Loss: 0.7881 | Learning Rate: 0.0009
Epoch 39 end
Epoch 40 start
[Epoch: 040] Contrastive Pre-train | Loss: 0.7916 | Learning Rate: 0.0009
Epoch 40 end
Epoch 41 start
[Epoch: 041] Contrastive Pre-train | Loss: 0.7725 | Learning Rate: 0.0008
Epoch 41 end
Epoch 42 start
[Epoch: 042] Contrastive Pre-train | Loss: 0.7688 | Learning Rate: 0.0008
Epoch 42 end
Epoch 43 start
[Epoch: 043] Contrastive Pre-train | Loss: 0.7546 | Learning Rate: 0.0008
Epoch 43 end
Epoch 44 start
[Epoch: 044] Contrastive Pre-train | Loss: 0.7666 | Learning Rate: 0.0008
Epoch 44 end
Epoch 45 start

[Epoch: 045] Contrastive Pre-train | Loss: 0.7491 | Learning Rate: 0.0008
Epoch 45 end
Epoch 46 start
[Epoch: 046] Contrastive Pre-train | Loss: 0.7613 | Learning Rate: 0.0008
Epoch 46 end
Epoch 47 start
[Epoch: 047] Contrastive Pre-train | Loss: 0.7541 | Learning Rate: 0.0008
Epoch 47 end
Epoch 48 start
[Epoch: 048] Contrastive Pre-train | Loss: 0.7554 | Learning Rate: 0.0008
Epoch 48 end
Epoch 49 start
[Epoch: 049] Contrastive Pre-train | Loss: 0.7467 | Learning Rate: 0.0008
Epoch 49 end
Epoch 50 start
[Epoch: 050] Contrastive Pre-train | Loss: 0.7570 | Learning Rate: 0.0008
Epoch 50 end

==> Saving...





```

Epoch 1 start
[Epoch: 001] Train      | Loss: 2.3026 | Accuracy: 0.7866
[Epoch: 001] Validation | Loss: 2.3025 | Accuracy: 0.8005
Epoch 1 end
Epoch 2 start
[Epoch: 002] Train      | Loss: 2.3026 | Accuracy: 0.7862
[Epoch: 002] Validation | Loss: 2.3025 | Accuracy: 0.8007
Epoch 2 end
> early stopping counter: 1 out of 30
Epoch 3 start
[Epoch: 003] Train      | Loss: 2.3026 | Accuracy: 0.7851
[Epoch: 003] Validation | Loss: 2.3025 | Accuracy: 0.7985
Epoch 3 end
> early stopping counter: 2 out of 30
Epoch 4 start
[Epoch: 004] Train      | Loss: 2.3026 | Accuracy: 0.7839
[Epoch: 004] Validation | Loss: 2.3025 | Accuracy: 0.8010
Epoch 4 end
> early stopping counter: 3 out of 30
Epoch 5 start

```

```
[Epoch: 005] Train      | Loss: 2.3026 | Accuracy: 0.7858
[Epoch: 005] Validation | Loss: 2.3025 | Accuracy: 0.8014
Epoch 5 end
> early stopping counter: 4 out of 30
Epoch 6 start
[Epoch: 006] Train      | Loss: 2.3026 | Accuracy: 0.7855
[Epoch: 006] Validation | Loss: 2.3025 | Accuracy: 0.7981
Epoch 6 end
> early stopping counter: 5 out of 30
Epoch 7 start
[Epoch: 007] Train      | Loss: 2.3026 | Accuracy: 0.7859
[Epoch: 007] Validation | Loss: 2.3025 | Accuracy: 0.7988
Epoch 7 end
> early stopping counter: 6 out of 30
Epoch 8 start
[Epoch: 008] Train      | Loss: 2.3026 | Accuracy: 0.7857
[Epoch: 008] Validation | Loss: 2.3025 | Accuracy: 0.8000
Epoch 8 end
> early stopping counter: 7 out of 30
Epoch 9 start
[Epoch: 009] Train      | Loss: 2.3026 | Accuracy: 0.7842
[Epoch: 009] Validation | Loss: 2.3025 | Accuracy: 0.8014
Epoch 9 end
> early stopping counter: 8 out of 30
Epoch 10 start
[Epoch: 010] Train      | Loss: 2.3026 | Accuracy: 0.7849
[Epoch: 010] Validation | Loss: 2.3025 | Accuracy: 0.7994
Epoch 10 end
> early stopping counter: 9 out of 30
Epoch 11 start
[Epoch: 011] Train      | Loss: 2.3026 | Accuracy: 0.7847
[Epoch: 011] Validation | Loss: 2.3025 | Accuracy: 0.8012
Epoch 11 end
> early stopping counter: 10 out of 30
Epoch 12 start
[Epoch: 012] Train      | Loss: 2.3026 | Accuracy: 0.7859
[Epoch: 012] Validation | Loss: 2.3025 | Accuracy: 0.8002
Epoch 12 end
> early stopping counter: 11 out of 30
Epoch 13 start
[Epoch: 013] Train      | Loss: 2.3026 | Accuracy: 0.7864
[Epoch: 013] Validation | Loss: 2.3025 | Accuracy: 0.7994
Epoch 13 end
> early stopping counter: 12 out of 30
Epoch 14 start
[Epoch: 014] Train      | Loss: 2.3026 | Accuracy: 0.7848
[Epoch: 014] Validation | Loss: 2.3025 | Accuracy: 0.8006
Epoch 14 end
```

```
> early stopping counter: 13 out of 30
Epoch 15 start
[Epoch: 015] Train      | Loss: 2.3026 | Accuracy: 0.7845
[Epoch: 015] Validation | Loss: 2.3025 | Accuracy: 0.7993
Epoch 15 end
> early stopping counter: 14 out of 30
Epoch 16 start
[Epoch: 016] Train      | Loss: 2.3026 | Accuracy: 0.7844
[Epoch: 016] Validation | Loss: 2.3025 | Accuracy: 0.7991
Epoch 16 end
> early stopping counter: 15 out of 30
Epoch 17 start
[Epoch: 017] Train      | Loss: 2.3026 | Accuracy: 0.7868
[Epoch: 017] Validation | Loss: 2.3025 | Accuracy: 0.7992
Epoch 17 end
> early stopping counter: 16 out of 30
Epoch 18 start
[Epoch: 018] Train      | Loss: 2.3026 | Accuracy: 0.7867
[Epoch: 018] Validation | Loss: 2.3025 | Accuracy: 0.8016
Epoch 18 end
> early stopping counter: 17 out of 30
Epoch 19 start
[Epoch: 019] Train      | Loss: 2.3026 | Accuracy: 0.7874
[Epoch: 019] Validation | Loss: 2.3025 | Accuracy: 0.7988
Epoch 19 end
> early stopping counter: 18 out of 30
Epoch 20 start
[Epoch: 020] Train      | Loss: 2.3026 | Accuracy: 0.7844
[Epoch: 020] Validation | Loss: 2.3025 | Accuracy: 0.7993
Epoch 20 end
> early stopping counter: 19 out of 30
Epoch 21 start
[Epoch: 021] Train      | Loss: 2.3026 | Accuracy: 0.7848
[Epoch: 021] Validation | Loss: 2.3025 | Accuracy: 0.8017
Epoch 21 end
> early stopping counter: 20 out of 30
Epoch 22 start
[Epoch: 022] Train      | Loss: 2.3026 | Accuracy: 0.7858
[Epoch: 022] Validation | Loss: 2.3025 | Accuracy: 0.8027
Epoch 22 end
> early stopping counter: 21 out of 30
Epoch 23 start
[Epoch: 023] Train      | Loss: 2.3026 | Accuracy: 0.7855
[Epoch: 023] Validation | Loss: 2.3025 | Accuracy: 0.8012
Epoch 23 end
> early stopping counter: 22 out of 30
Epoch 24 start
[Epoch: 024] Train      | Loss: 2.3026 | Accuracy: 0.7858
```

```
[Epoch: 024] Validation | Loss: 2.3025 | Accuracy: 0.7981
Epoch 24 end
> early stopping counter: 23 out of 30
Epoch 25 start
[Epoch: 025] Train      | Loss: 2.3026 | Accuracy: 0.7861
[Epoch: 025] Validation | Loss: 2.3025 | Accuracy: 0.7990
Epoch 25 end
> early stopping counter: 24 out of 30
Epoch 26 start
[Epoch: 026] Train      | Loss: 2.3026 | Accuracy: 0.7849
[Epoch: 026] Validation | Loss: 2.3025 | Accuracy: 0.8008
Epoch 26 end
> early stopping counter: 25 out of 30
Epoch 27 start
[Epoch: 027] Train      | Loss: 2.3026 | Accuracy: 0.7850
[Epoch: 027] Validation | Loss: 2.3025 | Accuracy: 0.7983
Epoch 27 end
> early stopping counter: 26 out of 30
Epoch 28 start
[Epoch: 028] Train      | Loss: 2.3026 | Accuracy: 0.7841
[Epoch: 028] Validation | Loss: 2.3025 | Accuracy: 0.8007
Epoch 28 end
> early stopping counter: 27 out of 30
Epoch 29 start
[Epoch: 029] Train      | Loss: 2.3026 | Accuracy: 0.7855
[Epoch: 029] Validation | Loss: 2.3025 | Accuracy: 0.7991
Epoch 29 end
> early stopping counter: 28 out of 30
Epoch 30 start
[Epoch: 030] Train      | Loss: 2.3026 | Accuracy: 0.7865
[Epoch: 030] Validation | Loss: 2.3025 | Accuracy: 0.8010
Epoch 30 end
> early stopping counter: 29 out of 30
Epoch 31 start
[Epoch: 031] Train      | Loss: 2.3026 | Accuracy: 0.7853
[Epoch: 031] Validation | Loss: 2.3025 | Accuracy: 0.8007
Epoch 31 end
> early stopping counter: 30 out of 30
Early stopping at epoch 31
```

